

```

from keras.utils import np_utils
from keras.datasets import mnist
import seaborn as sns
from keras.initializers import RandomNormal

```

↳ Using TensorFlow backend.

```

%matplotlib inline
import matplotlib.pyplot as plt
import numpy as np
import time
# https://gist.github.com/greydanus/f6eee59eaf1d90fcb3b534a25362cea4
# https://stackoverflow.com/a/14434334
# this function is used to update the plots for each epoch and error
def plt_dynamic(x, vy, ty, ax, colors=['b']):
    ax.plot(x, vy, 'b', label="Test Loss")
    ax.plot(x, ty, 'r', label="Train Loss")
    plt.legend()
    plt.grid()
    fig.canvas.draw()

```

```

# the data, shuffled and split between train and test sets
(X_train, y_train), (X_test, y_test) = mnist.load_data()

```

↳ Downloading data from <https://s3.amazonaws.com/img-datasets/mnist.npz>
11493376/11490434 [=====] - 0s 0us/step

```

print("Number of training examples :", X_train.shape[0], "and each image is of shape (%d, %d)"%(X_train.shape[1], X_train.shape[2]))
print("Number of testing examples :", X_test.shape[0], "and each image is of shape (%d, %d)"%(X_test.shape[1], X_test.shape[2]))

```

↳ Number of training examples : 60000 and each image is of shape (28, 28)
Number of testing examples : 10000 and each image is of shape (28, 28)

```

# if you observe the input shape its 2 dimensional vector
# for each image we have a (28*28) vector
# we will convert the (28*28) vector into single dimensional vector of 1 * 784

X_train = X_train.reshape(X_train.shape[0], X_train.shape[1]*X_train.shape[2])
X_test = X_test.reshape(X_test.shape[0], X_test.shape[1]*X_test.shape[2])

```

```

# after converting the input images from 3d to 2d vectors

```

```

print("Number of training examples :", X_train.shape[0], "and each image is of shape (%d)"%(X_train.shape[1]))
print("Number of testing examples :", X_test.shape[0], "and each image is of shape (%d)"%(X_test.shape[1]))

```

↳ Number of training examples : 60000 and each image is of shape (784)
Number of testing examples : 10000 and each image is of shape (784)

```

# An example data point
print(X_train[0])

```

↳

```
[
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  3  18  18  18 126 136 175 26 166 255
247 127  0  0  0  0  0  0  0  0  0  0  0  0  0  30 36 94 154
170 253 253 253 253 253 225 172 253 242 195 64  0  0  0  0  0  0
  0  0  0  0  0  49 238 253 253 253 253 253 253 253 251 93 82
 82 56 39  0  0  0  0  0  0  0  0  0  0  0  0 18 219 253
253 253 253 253 198 182 247 241  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0 80 156 107 253 253 205 11  0 43 154
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0 14  1 154 253 90  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0 139 253 190 2  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0 11 190 253 70  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0 35 241
225 160 108  1  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0 81 240 253 253 119 25  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0 45 186 253 253 150 27  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0 16 93 252 253 187
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0 249 253 249 64  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0 46 130 183 253
253 207  2  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0 39 148 229 253 253 253 250 182  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0 24 114 221 253 253 253
253 201 78  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0 23 66 213 253 253 253 253 198 81  2  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0 18 171 219 253 253 253 253 195
80  9  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
55 172 226 253 253 253 253 244 133 11  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0 136 253 253 253 212 135 132 16
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
```

```
# if we observe the above matrix each cell is having a value between 0-255
# before we move to apply machine learning algorithms lets try to normalize the data
# X => (X - Xmin)/(Xmax-Xmin) = X/255
```

```
X_train = X_train/255
X_test = X_test/255
```

```
# example data point after normlizing
print(X_train[0])
```



```
[0.      0.      0.      0.      0.      0.
 0.      0.      0.      0.      0.      0.
 0.      0.      0.      0.      0.      0.
 0.      0.      0.      0.      0.      0.
 0.      0.      0.      0.      0.      0.
 0.      0.      0.      0.      0.      0.
 0.      0.      0.      0.      0.      0.
 0.      0.      0.      0.      0.      0.
 0.      0.      0.      0.      0.      0.
 0.      0.      0.      0.      0.      0.
 0.      0.      0.      0.      0.      0.
 0.      0.      0.      0.      0.      0.
 0.      0.      0.      0.      0.      0.
 0.      0.      0.      0.      0.      0.
 0.      0.      0.      0.      0.      0.
 0.      0.      0.      0.      0.      0.
 0.      0.      0.      0.      0.      0.
 0.      0.      0.      0.      0.      0.
 0.      0.      0.      0.      0.      0.
 0.      0.      0.1176471 0.07058824 0.07058824 0.07058824
0.49411765 0.53333333 0.68627451 0.10196078 0.65098039 1.
0.96862745 0.49803922 0.      0.      0.      0.
0.      0.      0.      0.      0.      0.
0.      0.      0.11764706 0.14117647 0.36862745 0.60392157
0.66666667 0.99215686 0.99215686 0.99215686 0.99215686 0.99215686
0.88235294 0.6745098 0.99215686 0.94901961 0.76470588 0.25098039
0.      0.      0.      0.      0.      0.
0.      0.      0.      0.      0.      0.19215686
0.93333333 0.99215686 0.99215686 0.99215686 0.99215686 0.99215686
0.99215686 0.99215686 0.99215686 0.98431373 0.36470588 0.32156863
0.32156863 0.21960784 0.15294118 0.      0.      0.
0.      0.      0.      0.      0.      0.
0.      0.      0.      0.07058824 0.85882353 0.99215686
0.99215686 0.99215686 0.99215686 0.99215686 0.77647059 0.71372549
0.96862745 0.94509804 0.      0.      0.      0.
0.      0.      0.      0.      0.      0.
0.      0.      0.      0.      0.      0.
0.      0.      0.31372549 0.61176471 0.41960784 0.99215686
0.99215686 0.80392157 0.04313725 0.      0.16862745 0.60392157
0.      0.      0.      0.      0.      0.
0.      0.      0.      0.      0.      0.
0.      0.      0.      0.      0.      0.
0.      0.05490196 0.00392157 0.60392157 0.99215686 0.35294118
0.      0.      0.      0.      0.      0.
0.      0.      0.      0.      0.      0.
0.      0.      0.      0.      0.      0.
0.      0.      0.      0.      0.      0.
0.      0.      0.      0.      0.      0.
0.      0.54509804 0.99215686 0.74509804 0.00784314 0.
0.      0.      0.      0.      0.      0.
0.      0.      0.      0.      0.      0.
0.      0.      0.      0.      0.      0.
0.      0.      0.      0.      0.      0.04313725
0.74509804 0.99215686 0.2745098 0.      0.      0.
0.      0.      0.      0.      0.      0.
0.      0.      0.      0.      0.      0.]
```

<https://colab.research.google.com/drive/1mJH6PqJhh-v9xoEvzXcf8aGfVoPPjdez#scrollTo=pDc25elBmATb&printMode=true>

```

0. 0. 0. 0. 0. 0.
0. 0. 0. 0. 0. 0.
0. 0. 0. 0. 0. 0.
0. 0. 0. 0. 0. 0.

```

```

# here we are having a class number for each image
print("Class label of first image :", y_train[0])

# lets convert this into a 10 dimensional vector
# ex: consider an image is 5 convert it into 5 => [0, 0, 0, 0, 0, 1, 0, 0, 0, 0]
# this conversion needed for MLPs

Y_train = np_utils.to_categorical(y_train, 10)
Y_test = np_utils.to_categorical(y_test, 10)

print("After converting the output into a vector : ",Y_train[0])

```

```

[ ]> Class label of first image : 5
      After converting the output into a vector : [0. 0. 0. 0. 0. 1. 0. 0. 0. 0.]

```

```

# some model parameters

output_dim = 10
input_dim = X_train.shape[1]

batch_size = 128
nb_epoch = 20

```

MLP + Dropout + AdamOptimizer

1. Two hidden layers

```

# https://stackoverflow.com/questions/34716454/where-do-i-call-the-batchnormalization-function-in

from keras.models import Sequential
from keras.layers import Dense, Activation

from keras.layers import Dropout
from keras.layers.normalization import BatchNormalization

model = Sequential()

model.add(Dense(435, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal
model.add(BatchNormalization())
model.add(Dropout(0.5))

model.add(Dense(156, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, s
model.add(BatchNormalization())
model.add(Dropout(0.5))

model.add(Dense(output_dim, activation='softmax'))

model.summary()

```

```

[ ]>

```

WARNING: Logging before flag parsing goes to stderr.

W0825 20:34:50.848045 140160636204928 deprecation_wrapper.py:119] From /usr/local/lib

W0825 20:34:50.882638 140160636204928 deprecation_wrapper.py:119] From /usr/local/lib

W0825 20:34:50.890069 140160636204928 deprecation_wrapper.py:119] From /usr/local/lib

W0825 20:34:50.998440 140160636204928 deprecation_wrapper.py:119] From /usr/local/lib

W0825 20:34:51.022265 140160636204928 deprecation.py:506] From /usr/local/lib/python3

Instructions for updating:

Please use `rate` instead of `keep_prob`. Rate should be set to `rate = 1 - keep_prob`

W0825 20:34:51.133279 140160636204928 deprecation_wrapper.py:119] From /usr/local/lib

Layer (type)	Output Shape	Param #
dense_1 (Dense)	(None, 435)	341475
batch_normalization_1 (Batch Normalization)	(None, 435)	1740
dropout_1 (Dropout)	(None, 435)	0
dense_2 (Dense)	(None, 156)	68016
batch_normalization_2 (Batch Normalization)	(None, 156)	624
dropout_2 (Dropout)	(None, 156)	0
dense_3 (Dense)	(None, 10)	1570
Total params: 413,425		

```
model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
```

```
history = model.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_val, Y_val))
```



W0825 20:34:56.736942 140160636204928 deprecation_wrapper.py:119] From /usr/local/lib

W0825 20:34:56.864316 140160636204928 deprecation.py:323] From /usr/local/lib/python3

Instructions for updating:

Use tf.where in 2.0, which has the same broadcast rule as np.where

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 8s 135us/step - loss: 0.5083 - acc: 0.

Epoch 2/20

60000/60000 [=====] - 7s 118us/step - loss: 0.2395 - acc: 0.

Epoch 3/20

60000/60000 [=====] - 7s 114us/step - loss: 0.1836 - acc: 0.

Epoch 4/20

60000/60000 [=====] - 7s 115us/step - loss: 0.1547 - acc: 0.

Epoch 5/20

60000/60000 [=====] - 7s 114us/step - loss: 0.1381 - acc: 0.

Epoch 6/20

60000/60000 [=====] - 7s 114us/step - loss: 0.1196 - acc: 0.

Epoch 7/20

60000/60000 [=====] - 7s 117us/step - loss: 0.1133 - acc: 0.

Epoch 8/20

60000/60000 [=====] - 7s 116us/step - loss: 0.1042 - acc: 0.

Epoch 9/20

60000/60000 [=====] - 7s 116us/step - loss: 0.0970 - acc: 0.

Epoch 10/20

60000/60000 [=====] - 7s 117us/step - loss: 0.0894 - acc: 0.

Epoch 11/20

60000/60000 [=====] - 7s 112us/step - loss: 0.0855 - acc: 0.

Epoch 12/20

60000/60000 [=====] - 7s 115us/step - loss: 0.0813 - acc: 0.

Epoch 13/20

60000/60000 [=====] - 7s 115us/step - loss: 0.0765 - acc: 0.

Epoch 14/20

60000/60000 [=====] - 7s 114us/step - loss: 0.0712 - acc: 0.

Epoch 15/20

60000/60000 [=====] - 7s 116us/step - loss: 0.0713 - acc: 0.

Epoch 16/20

60000/60000 [=====] - 7s 114us/step - loss: 0.0677 - acc: 0.

Epoch 17/20

60000/60000 [=====] - 7s 114us/step - loss: 0.0654 - acc: 0.

Epoch 18/20

60000/60000 [=====] - 7s 116us/step - loss: 0.0604 - acc: 0.

Epoch 19/20

60000/60000 [=====] - 7s 114us/step - loss: 0.0554 - acc: 0.

score = model.evaluate(X_test, Y_test, verbose=0)

print('Test score:', score[0])

print('Test accuracy:', score[1])

fig,ax = plt.subplots(1,1)

ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')

list of epoch numbers

x = list(range(1,nb_epoch+1))

val_loss : validation loss

val_acc : validation accuracy

loss : training loss


acc : train accuracy

vy = history.history['val_loss']

ty = history.history['loss']

plt_dynamic(x, vy, ty, ax)

Test score: 0.054573897721024694
 Test accuracy: 0.9833




```
w_after = model.get_weights()

h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)

fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')

plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



2. Three hidden layers

```
# https://stackoverflow.com/questions/34716454/where-do-i-call-the-batchnormalization-function-in

from keras.models import Sequential
from keras.layers import Dense, Activation

from keras.layers import Dropout
from keras.layers.normalization import BatchNormalization

model = Sequential()

model.add(Dense(506, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal
model.add(BatchNormalization())
model.add(Dropout(0.5))

model.add(Dense(290, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, s
model.add(BatchNormalization())
model.add(Dropout(0.5))

model.add(Dense(130, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.225, s
model.add(BatchNormalization())
model.add(Dropout(0.5))

model.add(Dense(output_dim, activation='softmax'))

model.summary()
```



Layer (type)	Output Shape	Param #
dense_4 (Dense)	(None, 506)	397210
batch_normalization_3 (Batch Normalization)	(None, 506)	2024
dropout_3 (Dropout)	(None, 506)	0
dense_5 (Dense)	(None, 290)	147030
batch_normalization_4 (Batch Normalization)	(None, 290)	1160
dropout_4 (Dropout)	(None, 290)	0
dense_6 (Dense)	(None, 130)	37830
batch_normalization_5 (Batch Normalization)	(None, 130)	520
dropout_5 (Dropout)	(None, 130)	0
dense_7 (Dense)	(None, 10)	1310

```
model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
history = model.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validation_data=(X_val, Y_val))
```



Train on 60000 samples, validate on 10000 samples

```
Epoch 1/20
60000/60000 [=====] - 11s 187us/step - loss: 0.6955 - acc: 0
Epoch 2/20
60000/60000 [=====] - 10s 173us/step - loss: 0.2897 - acc: 0
Epoch 3/20
60000/60000 [=====] - 10s 173us/step - loss: 0.2240 - acc: 0
Epoch 4/20
60000/60000 [=====] - 10s 173us/step - loss: 0.1812 - acc: 0
Epoch 5/20
60000/60000 [=====] - 10s 171us/step - loss: 0.1579 - acc: 0
Epoch 6/20
60000/60000 [=====] - 10s 174us/step - loss: 0.1461 - acc: 0
Epoch 7/20
60000/60000 [=====] - 10s 170us/step - loss: 0.1338 - acc: 0
Epoch 8/20
60000/60000 [=====] - 10s 172us/step - loss: 0.1180 - acc: 0
Epoch 9/20
60000/60000 [=====] - 10s 172us/step - loss: 0.1122 - acc: 0
Epoch 10/20
60000/60000 [=====] - 10s 170us/step - loss: 0.1032 - acc: 0
Epoch 11/20
60000/60000 [=====] - 10s 173us/step - loss: 0.0964 - acc: 0
Epoch 12/20
60000/60000 [=====] - 10s 174us/step - loss: 0.0915 - acc: 0
Epoch 13/20
60000/60000 [=====] - 10s 173us/step - loss: 0.0845 - acc: 0
Epoch 14/20
60000/60000 [=====] - 10s 172us/step - loss: 0.0826 - acc: 0
Epoch 15/20
60000/60000 [=====] - 10s 173us/step - loss: 0.0796 - acc: 0
Epoch 16/20
60000/60000 [=====] - 10s 172us/step - loss: 0.0776 - acc: 0
Epoch 17/20
60000/60000 [=====] - 11s 175us/step - loss: 0.0771 - acc: 0
Epoch 18/20
60000/60000 [=====] - 10s 174us/step - loss: 0.0697 - acc: 0
```

```
score = model.evaluate(X_test, Y_test, verbose=0)
```

```
print('Test score:', score[0])
```

```
print('Test accuracy:', score[1])
```

```
fig,ax = plt.subplots(1,1)
```

```
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')
```

```
# list of epoch numbers
```

```
x = list(range(1,nb_epoch+1))
```

```
# val_loss : validation loss
```

```
# val_acc : validation accuracy
```

```
# loss : training loss
```

```
# acc : train accuracy
```

```
vy = history.history['val_loss']
```


```
ty = history.history['loss']
```

```
plt_dynamic(x, vy, ty, ax)
```



Test score: 0.06090008487093437

Test accuracy: 0.9841




```
w_after = model.get_weights()
```

```
h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)
```

```
fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')
```

```
plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')
```

```
plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



3. Five hidden layers

```
model = Sequential()
```

```
model.add(Dense(609, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal))
model.add(BatchNormalization())
model.add(Dropout(0.5))
```

```

model.add(Dense(510, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.125, s
model.add(BatchNormalization())
model.add(Dropout(0.5))

model.add(Dense(390, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.205, s
model.add(BatchNormalization())
model.add(Dropout(0.5))

model.add(Dense(270, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.285, s
model.add(BatchNormalization())
model.add(Dropout(0.5))

model.add(Dense(110, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.325, s
model.add(BatchNormalization())
model.add(Dropout(0.5))

model.add(Dense(output_dim, activation='softmax'))

model.summary()

```



Layer (type)	Output Shape	Param #
=====		
dense_8 (Dense)	(None, 609)	478065
batch_normalization_6 (Batch Normalization)	(None, 609)	2436
dropout_6 (Dropout)	(None, 609)	0
dense_9 (Dense)	(None, 510)	311100
batch_normalization_7 (Batch Normalization)	(None, 510)	2040
dropout_7 (Dropout)	(None, 510)	0
dense_10 (Dense)	(None, 390)	199290
batch_normalization_8 (Batch Normalization)	(None, 390)	1560
dropout_8 (Dropout)	(None, 390)	0
dense_11 (Dense)	(None, 270)	105570
batch_normalization_9 (Batch Normalization)	(None, 270)	1080
dropout_9 (Dropout)	(None, 270)	0
dense_12 (Dense)	(None, 110)	29810
batch_normalization_10 (Batch Normalization)	(None, 110)	440
dropout_10 (Dropout)	(None, 110)	0
dense_13 (Dense)	(None, 10)	1110
=====		
Total params: 1,132,501		
Trainable params: 1,128,723		
Non-trainable params: 3,778		

```
model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])  
history = model.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, validati
```



Train on 60000 samples, validate on 10000 samples

```
Epoch 1/20
60000/60000 [=====] - 21s 358us/step - loss: 1.3110 - acc: 0
Epoch 2/20
60000/60000 [=====] - 20s 326us/step - loss: 0.4553 - acc: 0
Epoch 3/20
60000/60000 [=====] - 20s 328us/step - loss: 0.3142 - acc: 0
Epoch 4/20
60000/60000 [=====] - 20s 327us/step - loss: 0.2511 - acc: 0
Epoch 5/20
60000/60000 [=====] - 20s 327us/step - loss: 0.2115 - acc: 0
Epoch 6/20
60000/60000 [=====] - 20s 327us/step - loss: 0.1863 - acc: 0
Epoch 7/20
60000/60000 [=====] - 20s 327us/step - loss: 0.1643 - acc: 0
Epoch 8/20
60000/60000 [=====] - 20s 327us/step - loss: 0.1540 - acc: 0
Epoch 9/20
60000/60000 [=====] - 19s 324us/step - loss: 0.1399 - acc: 0
Epoch 10/20
60000/60000 [=====] - 20s 328us/step - loss: 0.1322 - acc: 0
Epoch 11/20
60000/60000 [=====] - 20s 326us/step - loss: 0.1216 - acc: 0
Epoch 12/20
60000/60000 [=====] - 20s 332us/step - loss: 0.1200 - acc: 0
Epoch 13/20
60000/60000 [=====] - 20s 329us/step - loss: 0.1092 - acc: 0
Epoch 14/20
60000/60000 [=====] - 20s 329us/step - loss: 0.1043 - acc: 0
Epoch 15/20
60000/60000 [=====] - 19s 322us/step - loss: 0.1009 - acc: 0
Epoch 16/20
60000/60000 [=====] - 19s 325us/step - loss: 0.0983 - acc: 0
Epoch 17/20
60000/60000 [=====] - 20s 327us/step - loss: 0.0936 - acc: 0
Epoch 18/20
60000/60000 [=====] - 20s 326us/step - loss: 0.0882 - acc: 0
```

```
score = model.evaluate(X_test, Y_test, verbose=0)
```

```
print('Test score:', score[0])
```

```
print('Test accuracy:', score[1])
```

```
fig,ax = plt.subplots(1,1)
```

```
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')
```

```
# list of epoch numbers
```

```
x = list(range(1,nb_epoch+1))
```

```
# val_loss : validation loss
```

```
# val_acc : validation accuracy
```

```
# loss : training loss
```

```
# acc : train accuracy
```

```
vy = history.history['val_loss']
```


```
ty = history.history['loss']
```

```
plt_dynamic(x, vy, ty, ax)
```



Test score: 0.07164842545758002

Test accuracy: 0.9812




```
w_after = model.get_weights()
```

```
h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)
```

```
fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')
```

```
plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')
```

```
plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



MLP + Sigmoid activation + SGDOptimizer

```
#three hidden layers
model_sigmoid = Sequential()
model_sigmoid.add(Dense(512, activation='sigmoid', input_shape=(input_dim,)))
model_sigmoid.add(Dense(128, activation='sigmoid'))
model_sigmoid.add(Dense(65, activation='sigmoid'))
```



```
model_sigmoid.add(Dense(output_dim, activation='softmax'))

model_sigmoid.summary()
```

↗

Layer (type)	Output Shape	Param #
=====		
dense_14 (Dense)	(None, 512)	401920

dense_15 (Dense)	(None, 128)	65664

dense_16 (Dense)	(None, 65)	8385

dense_17 (Dense)	(None, 10)	660
=====		
Total params: 476,629		
Trainable params: 476,629		
Non-trainable params: 0		

```
model_sigmoid.compile(optimizer='sgd', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_sigmoid.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1,
```



Train on 60000 samples, validate on 10000 samples

```
Epoch 1/20
60000/60000 [=====] - 6s 92us/step - loss: 2.3097 - acc: 0.1
Epoch 2/20
60000/60000 [=====] - 5s 83us/step - loss: 2.2949 - acc: 0.1
Epoch 3/20
60000/60000 [=====] - 5s 83us/step - loss: 2.2887 - acc: 0.1
Epoch 4/20
60000/60000 [=====] - 5s 81us/step - loss: 2.2820 - acc: 0.1
Epoch 5/20
60000/60000 [=====] - 5s 82us/step - loss: 2.2742 - acc: 0.1
Epoch 6/20
60000/60000 [=====] - 5s 83us/step - loss: 2.2646 - acc: 0.1
Epoch 7/20
60000/60000 [=====] - 5s 80us/step - loss: 2.2525 - acc: 0.2
Epoch 8/20
60000/60000 [=====] - 5s 85us/step - loss: 2.2368 - acc: 0.2
Epoch 9/20
60000/60000 [=====] - 5s 85us/step - loss: 2.2157 - acc: 0.3
Epoch 10/20
60000/60000 [=====] - 5s 79us/step - loss: 2.1858 - acc: 0.4
Epoch 11/20
60000/60000 [=====] - 5s 82us/step - loss: 2.1419 - acc: 0.4
Epoch 12/20
60000/60000 [=====] - 5s 84us/step - loss: 2.0762 - acc: 0.5
Epoch 13/20
60000/60000 [=====] - 5s 86us/step - loss: 1.9797 - acc: 0.5
Epoch 14/20
60000/60000 [=====] - 5s 85us/step - loss: 1.8517 - acc: 0.5
Epoch 15/20
60000/60000 [=====] - 5s 83us/step - loss: 1.7101 - acc: 0.5
Epoch 16/20
60000/60000 [=====] - 5s 84us/step - loss: 1.5782 - acc: 0.5
Epoch 17/20
60000/60000 [=====] - 5s 83us/step - loss: 1.4647 - acc: 0.6
Epoch 18/20
60000/60000 [=====] - 5s 81us/step - loss: 1.3658 - acc: 0.6
```

```
score = model_sigmoid.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])
```

```
fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')
```


```
# list of epoch numbers
x = list(range(1,nb_epoch+1)).
```

```
vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```



Test score: 1.1352740411758422

Test accuracy: 0.6817




```
w_after = model_sigmoid.get_weights()
```

```
h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)
```

```
fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')
```

```
plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')
```

```
plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



MLP + Sigmoid activation + ADAM

```
#five hidden layers with BN and Dropout
model_sigmoid = Sequential()
```

```
model_sigmoid.add(Dense(665, activation='sigmoid', input_shape=(input_dim,)))
model_sigmoid.add(BatchNormalization())
```

```
model_sigmoid.add(Dropout(0.2))

model_sigmoid.add(Dense(456, activation='sigmoid'))
model_sigmoid.add(BatchNormalization())
model_sigmoid.add(Dropout(0.2))

model_sigmoid.add(Dense(348, activation='sigmoid'))
model_sigmoid.add(BatchNormalization())
model_sigmoid.add(Dropout(0.2))

model_sigmoid.add(Dense(250, activation='sigmoid'))
model_sigmoid.add(BatchNormalization())
model_sigmoid.add(Dropout(0.2))

model_sigmoid.add(Dense(115, activation='sigmoid'))
model_sigmoid.add(BatchNormalization())
model_sigmoid.add(Dropout(0.2))

model_sigmoid.add(Dense(output_dim, activation='softmax'))

model_sigmoid.summary()

model_sigmoid.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_sigmoid.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1,
```



Layer (type)	Output Shape	Param #
dense_47 (Dense)	(None, 665)	522025
batch_normalization_30 (Batch Normalization)	(None, 665)	2660
dropout_25 (Dropout)	(None, 665)	0
dense_48 (Dense)	(None, 456)	303696
batch_normalization_31 (Batch Normalization)	(None, 456)	1824
dropout_26 (Dropout)	(None, 456)	0
dense_49 (Dense)	(None, 348)	159036
batch_normalization_32 (Batch Normalization)	(None, 348)	1392
dropout_27 (Dropout)	(None, 348)	0
dense_50 (Dense)	(None, 250)	87250
batch_normalization_33 (Batch Normalization)	(None, 250)	1000
dropout_28 (Dropout)	(None, 250)	0
dense_51 (Dense)	(None, 115)	28865
batch_normalization_34 (Batch Normalization)	(None, 115)	460
dropout_29 (Dropout)	(None, 115)	0
dense_52 (Dense)	(None, 10)	1160
Total params: 1,109,368		
Trainable params: 1,105,700		
Non-trainable params: 3,668		

Train on 60000 samples, validate on 10000 samples

Epoch 1/20

60000/60000 [=====] - 22s 362us/step - loss: 0.3562 - acc: 0

Epoch 2/20

60000/60000 [=====] - 18s 307us/step - loss: 0.2023 - acc: 0

Epoch 3/20

60000/60000 [=====] - 19s 309us/step - loss: 0.1596 - acc: 0

Epoch 4/20

60000/60000 [=====] - 19s 309us/step - loss: 0.1305 - acc: 0

Epoch 5/20

60000/60000 [=====] - 19s 310us/step - loss: 0.1109 - acc: 0

Epoch 6/20

60000/60000 [=====] - 19s 313us/step - loss: 0.1000 - acc: 0

Epoch 7/20

60000/60000 [=====] - 19s 312us/step - loss: 0.0868 - acc: 0

Epoch 8/20

60000/60000 [=====] - 19s 311us/step - loss: 0.0791 - acc: 0

Epoch 9/20

60000/60000 [=====] - 19s 311us/step - loss: 0.0694 - acc: 0

Epoch 10/20

60000/60000 [=====] - 19s 311us/step - loss: 0.0661 - acc: 0

Epoch 11/20

```

60000/60000 [=====] - 19s 312us/step - loss: 0.0580 - acc: 0
Epoch 12/20
60000/60000 [=====] - 19s 310us/step - loss: 0.0560 - acc: 0
Epoch 13/20
60000/60000 [=====] - 19s 312us/step - loss: 0.0515 - acc: 0
Epoch 14/20
60000/60000 [=====] - 19s 310us/step - loss: 0.0471 - acc: 0
Epoch 15/20
60000/60000 [=====] - 19s 309us/step - loss: 0.0417 - acc: 0
Epoch 16/20
60000/60000 [=====] - 19s 313us/step - loss: 0.0437 - acc: 0
Epoch 17/20
60000/60000 [=====] - 19s 308us/step - loss: 0.0378 - acc: 0
Epoch 18/20
60000/60000 [=====] - 19s 310us/step - loss: 0.0367 - acc: 0

```

```
score = model_sigmoid.evaluate(X_test, Y_test, verbose=0)
```

```
print('Test score:', score[0])
```

```
print('Test accuracy:', score[1])
```

```
fig, ax = plt.subplots(1, 1)
```

```
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')
```

```
# list of epoch numbers
```

```
x = list(range(1, nb_epoch+1))
```

```
vy = history.history['val_loss']
```


```
ty = history.history['loss']
```

```
plt_dynamic(x, vy, ty, ax)
```

```

↳ Test score: 0.0672482188842725
Test accuracy: 0.9833

```



```
w_after = model_sigmoid.get_weights()
```

```
h1_w = w_after[0].flatten().reshape(-1,1)
```

```
h2_w = w_after[2].flatten().reshape(-1,1)
```

```
out_w = w_after[4].flatten().reshape(-1,1)
```

```
fig = plt.figure()
```

```
plt.title("Weight matrices after model trained")
```

```
plt.subplot(1, 3, 1)
```

```
plt.title("Trained model Weights")
```

```
ax = sns.violinplot(y=h1_w, color='b')
```

```
plt.xlabel('Hidden Layer 1')
```


```
plt.subplot(1, 3, 2)
```

```
plt.title("Trained model Weights")
```

```
ax = sns.violinplot(y=h2_w, color='r')
```

```
plt.xlabel('Hidden Layer 2 ')

plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



MLP + ReLU +SGD

#three hidden layers with different weight initializations

```
model_relu = Sequential()
model_relu.add(Dense(548, activation='relu', input_shape=(input_dim,), kernel_initializer=RandomNormal(mean=0.0, stddev=0.1))
model_relu.add(Dense(328, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.1))
model_relu.add(Dense(158, activation='relu', kernel_initializer=RandomNormal(mean=0.0, stddev=0.1))
model_relu.add(Dense(output_dim, activation='softmax'))
```

```
model_relu.summary()
```

Layer (type)	Output Shape	Param #
dense_25 (Dense)	(None, 548)	430180
dense_26 (Dense)	(None, 328)	180072
dense_27 (Dense)	(None, 158)	51982
dense_28 (Dense)	(None, 10)	1590
Total params: 663,824		
Trainable params: 663,824		
Non-trainable params: 0		

```
model_relu.compile(optimizer='sgd', loss='categorical_crossentropy', metrics=['accuracy'])
history = model_relu.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, val
```

↳

Train on 60000 samples, validate on 10000 samples

```
Epoch 1/20
60000/60000 [=====] - 8s 137us/step - loss: 0.4969 - acc: 0.
Epoch 2/20
60000/60000 [=====] - 7s 124us/step - loss: 0.2506 - acc: 0.
Epoch 3/20
60000/60000 [=====] - 7s 118us/step - loss: 0.1990 - acc: 0.
Epoch 4/20
60000/60000 [=====] - 7s 122us/step - loss: 0.1674 - acc: 0.
Epoch 5/20
60000/60000 [=====] - 7s 118us/step - loss: 0.1456 - acc: 0.
Epoch 6/20
60000/60000 [=====] - 7s 121us/step - loss: 0.1287 - acc: 0.
Epoch 7/20
60000/60000 [=====] - 7s 122us/step - loss: 0.1160 - acc: 0.
Epoch 8/20
60000/60000 [=====] - 7s 122us/step - loss: 0.1049 - acc: 0.
Epoch 9/20
60000/60000 [=====] - 7s 122us/step - loss: 0.0958 - acc: 0.
Epoch 10/20
60000/60000 [=====] - 7s 119us/step - loss: 0.0880 - acc: 0.
Epoch 11/20
60000/60000 [=====] - 7s 121us/step - loss: 0.0809 - acc: 0.
Epoch 12/20
60000/60000 [=====] - 7s 120us/step - loss: 0.0750 - acc: 0.
Epoch 13/20
60000/60000 [=====] - 7s 122us/step - loss: 0.0694 - acc: 0.
Epoch 14/20
60000/60000 [=====] - 7s 120us/step - loss: 0.0648 - acc: 0.
Epoch 15/20
60000/60000 [=====] - 7s 117us/step - loss: 0.0603 - acc: 0.
Epoch 16/20
60000/60000 [=====] - 7s 117us/step - loss: 0.0562 - acc: 0.
Epoch 17/20
60000/60000 [=====] - 7s 119us/step - loss: 0.0526 - acc: 0.
Epoch 18/20
60000/60000 [=====] - 7s 121us/step - loss: 0.0493 - acc: 0.
```

```
score = model_relu.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])
```

```
fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')
```


```
# list of epoch numbers
x = list(range(1,nb_epoch+1))
```

```
vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```



Test score: 0.10271617978247814

Test accuracy: 0.9711




```
w_after = model_relu.get_weights()
```

```
h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)
```

```
fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')
```

```
plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')
```

```
plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



MLP + Batch-Norm on hidden Layers + AdamOptimizer

```
#Five hidden layers with activation='sigmoid', BN & different weight initializations
```

```
model_batch = Sequential()
```

```
model_batch.add(Dense(512, activation='sigmoid', input_shape=(input_dim,), kernel_initializer=RandomNormal(0.01, 0.01)))
```

```

model_batch.add(BatchNormalization())

model_batch.add(Dense(456, activation='sigmoid', kernel_initializer=RandomNormal(mean=0.0, stddev
model_batch.add(BatchNormalization())

model_batch.add(Dense(368, activation='sigmoid', kernel_initializer=RandomNormal(mean=0.0, stddev
model_batch.add(BatchNormalization())

model_batch.add(Dense(228, activation='sigmoid', kernel_initializer=RandomNormal(mean=0.0, stddev
model_batch.add(BatchNormalization())

model_batch.add(Dense(123, activation='sigmoid', kernel_initializer=RandomNormal(mean=0.0, stddev
model_batch.add(BatchNormalization())

model_batch.add(Dense(output_dim, activation='softmax'))

model_batch.summary()

```



Layer (type)	Output Shape	Param #
=====	=====	=====
dense_33 (Dense)	(None, 512)	401920
batch_normalization_19 (Batch Normalization)	(None, 512)	2048
dense_34 (Dense)	(None, 456)	233928
batch_normalization_20 (Batch Normalization)	(None, 456)	1824
dense_35 (Dense)	(None, 368)	168176
batch_normalization_21 (Batch Normalization)	(None, 368)	1472
dense_36 (Dense)	(None, 228)	84132
batch_normalization_22 (Batch Normalization)	(None, 228)	912
dense_37 (Dense)	(None, 123)	28167
batch_normalization_23 (Batch Normalization)	(None, 123)	492
dense_38 (Dense)	(None, 10)	1240
=====	=====	=====
Total params: 924,311		
Trainable params: 920,937		
Non-trainable params: 3,374		
=====	=====	=====

```

model_batch.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_batch.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, va

```



Train on 60000 samples, validate on 10000 samples

```
Epoch 1/20
60000/60000 [=====] - 18s 294us/step - loss: 0.3484 - acc: 0
Epoch 2/20
60000/60000 [=====] - 15s 249us/step - loss: 0.1872 - acc: 0
Epoch 3/20
60000/60000 [=====] - 15s 252us/step - loss: 0.1451 - acc: 0
Epoch 4/20
60000/60000 [=====] - 15s 252us/step - loss: 0.1237 - acc: 0
Epoch 5/20
60000/60000 [=====] - 15s 252us/step - loss: 0.1047 - acc: 0
Epoch 6/20
60000/60000 [=====] - 15s 248us/step - loss: 0.0951 - acc: 0
Epoch 7/20
60000/60000 [=====] - 15s 253us/step - loss: 0.0879 - acc: 0
Epoch 8/20
60000/60000 [=====] - 15s 253us/step - loss: 0.0779 - acc: 0
Epoch 9/20
60000/60000 [=====] - 15s 250us/step - loss: 0.0735 - acc: 0
Epoch 10/20
60000/60000 [=====] - 15s 250us/step - loss: 0.0671 - acc: 0
Epoch 11/20
60000/60000 [=====] - 15s 248us/step - loss: 0.0617 - acc: 0
Epoch 12/20
60000/60000 [=====] - 15s 245us/step - loss: 0.0560 - acc: 0
Epoch 13/20
60000/60000 [=====] - 15s 248us/step - loss: 0.0557 - acc: 0
Epoch 14/20
60000/60000 [=====] - 15s 252us/step - loss: 0.0529 - acc: 0
Epoch 15/20
60000/60000 [=====] - 15s 250us/step - loss: 0.0510 - acc: 0
Epoch 16/20
60000/60000 [=====] - 15s 249us/step - loss: 0.0467 - acc: 0
Epoch 17/20
60000/60000 [=====] - 15s 248us/step - loss: 0.0451 - acc: 0
Epoch 18/20
60000/60000 [=====] - 15s 252us/step - loss: 0.0416 - acc: 0
```

```
score = model_batch.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])
```

```
fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')
```


```
# list of epoch numbers
x = list(range(1,nb_epoch+1))
```

```
vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```



Test score: 0.10247163554290309

Test accuracy: 0.9707




```
w_after = model_batch.get_weights()
```

```
h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)
```

```
fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')
```

```
plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')
```

```
plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



MLP + Dropout + AdamOptimizer

```
#three hidden layers with different weight initializations
model_drop = Sequential()
```

```
model_drop.add(Dense(412, activation='sigmoid', input_shape=(input_dim,)), kernel_initializer=Rand
model_drop.add(BatchNormalization())
```

```

model_drop.add(Dropout(0.3))

model_drop.add(Dense(228, activation='sigmoid', kernel_initializer=RandomNormal(mean=0.0, stddev=0.01)))
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.3))

model_drop.add(Dense(98, activation='sigmoid', kernel_initializer=RandomNormal(mean=0.0, stddev=0.01)))
model_drop.add(BatchNormalization())
model_drop.add(Dropout(0.3))

model_drop.add(Dense(output_dim, activation='softmax'))

model_drop.summary()

```



Layer (type)	Output Shape	Param #
=====		
dense_43 (Dense)	(None, 412)	323420
batch_normalization_27 (Batch Normalization)	(None, 412)	1648
dropout_22 (Dropout)	(None, 412)	0
dense_44 (Dense)	(None, 228)	94164
batch_normalization_28 (Batch Normalization)	(None, 228)	912
dropout_23 (Dropout)	(None, 228)	0
dense_45 (Dense)	(None, 98)	22442
batch_normalization_29 (Batch Normalization)	(None, 98)	392
dropout_24 (Dropout)	(None, 98)	0
dense_46 (Dense)	(None, 10)	990
=====		
Total params: 443,968		
Trainable params: 442,492		
Non-trainable params: 1,476		
=====		

```

model_drop.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

history = model_drop.fit(X_train, Y_train, batch_size=batch_size, epochs=nb_epoch, verbose=1, val

```



Train on 60000 samples, validate on 10000 samples

```
Epoch 1/20
60000/60000 [=====] - 11s 178us/step - loss: 0.7992 - acc: 0.
Epoch 2/20
60000/60000 [=====] - 8s 133us/step - loss: 0.4700 - acc: 0.
Epoch 3/20
60000/60000 [=====] - 8s 131us/step - loss: 0.4104 - acc: 0.
Epoch 4/20
60000/60000 [=====] - 8s 133us/step - loss: 0.3763 - acc: 0.
Epoch 5/20
60000/60000 [=====] - 8s 133us/step - loss: 0.3557 - acc: 0.
Epoch 6/20
60000/60000 [=====] - 8s 134us/step - loss: 0.3318 - acc: 0.
Epoch 7/20
60000/60000 [=====] - 8s 132us/step - loss: 0.3157 - acc: 0.
Epoch 8/20
60000/60000 [=====] - 8s 133us/step - loss: 0.3036 - acc: 0.
Epoch 9/20
60000/60000 [=====] - 8s 134us/step - loss: 0.2891 - acc: 0.
Epoch 10/20
60000/60000 [=====] - 8s 133us/step - loss: 0.2730 - acc: 0.
Epoch 11/20
60000/60000 [=====] - 8s 133us/step - loss: 0.2674 - acc: 0.
Epoch 12/20
60000/60000 [=====] - 8s 138us/step - loss: 0.2542 - acc: 0.
Epoch 13/20
60000/60000 [=====] - 8s 136us/step - loss: 0.2451 - acc: 0.
Epoch 14/20
60000/60000 [=====] - 8s 137us/step - loss: 0.2316 - acc: 0.
Epoch 15/20
60000/60000 [=====] - 8s 136us/step - loss: 0.2220 - acc: 0.
Epoch 16/20
60000/60000 [=====] - 8s 132us/step - loss: 0.2155 - acc: 0.
Epoch 17/20
60000/60000 [=====] - 8s 135us/step - loss: 0.2024 - acc: 0.
Epoch 18/20
60000/60000 [=====] - 8s 133us/step - loss: 0.1932 - acc: 0.
```

```
score = model_drop.evaluate(X_test, Y_test, verbose=0)
print('Test score:', score[0])
print('Test accuracy:', score[1])
```

```
fig,ax = plt.subplots(1,1)
ax.set_xlabel('epoch') ; ax.set_ylabel('Categorical Crossentropy Loss')
```


```
# list of epoch numbers
x = list(range(1,nb_epoch+1))
```

```
vy = history.history['val_loss']
ty = history.history['loss']
plt_dynamic(x, vy, ty, ax)
```



Test score: 0.12330705345384776

Test accuracy: 0.9624




```
w_after = model_drop.get_weights()
```

```
h1_w = w_after[0].flatten().reshape(-1,1)
h2_w = w_after[2].flatten().reshape(-1,1)
out_w = w_after[4].flatten().reshape(-1,1)
```

```
fig = plt.figure()
plt.title("Weight matrices after model trained")
plt.subplot(1, 3, 1)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h1_w,color='b')
plt.xlabel('Hidden Layer 1')
```

```
plt.subplot(1, 3, 2)
plt.title("Trained model Weights")
ax = sns.violinplot(y=h2_w, color='r')
plt.xlabel('Hidden Layer 2 ')
```

```
plt.subplot(1, 3, 3)
plt.title("Trained model Weights")
ax = sns.violinplot(y=out_w,color='y')
plt.xlabel('Output Layer ')
plt.show()
```



Conclusion.

```
from prettytable import PrettyTable
```

```
x = PrettyTable()
```

```
x.field_names = ["Activation Function", "optimizer", "Batch Normalisations", "Dropout", "No. of hidde
```

```

x.add_row(["relu", "AdamOptimizer", "Yes", 0.5, 2, 0.054, 0.98])
x.add_row(["relu", "AdamOptimizer", "Yes", 0.5, 3, 0.060, 0.98])
x.add_row(["relu", "AdamOptimizer", "Yes", 0.5, 5, 0.071, 0.98])
x.add_row(["relu", "SGDOptimizer", "No", "No dropout", 3, 0.102, 0.97])
x.add_row(["Sigmoid", "SGDOptimizer", "No", "No dropout", 3, 1.13, 0.68])
x.add_row(["Sigmoid", "AdamOptimizer", "Yes", 0.2, 5, 0.067, 0.98])
x.add_row(["Sigmoid", "AdamOptimizer", "Yes", "No dropout", 5, 0.102, 0.97])
x.add_row(["Sigmoid", "AdamOptimizer", "Yes", 0.3, 3, 0.123, 0.96])
print(x)

```

↳

Activation Function	optimizer	Batch Normalisations	Dropout	No. of hi
relu	AdamOptimizer	Yes	0.5	
relu	AdamOptimizer	Yes	0.5	
relu	AdamOptimizer	Yes	0.5	
relu	SGDOptimizer	No	No dropout	
Sigmoid	SGDOptimizer	No	No dropout	
Sigmoid	AdamOptimizer	Yes	0.2	
Sigmoid	AdamOptimizer	Yes	No dropout	
Sigmoid	AdamOptimizer	Yes	0.3	

Observations:

1. We can see with different architecture, activation function & optimizer test score and test acc
2. from above table, we can observe relu activation have minimum test score and maximum tes
3. AdamOptimizer is better than sgdoptimizer
4. Batch Normalisations and Dropout are playing important role in improving test accuracy
5. MLP having sigmoid activation and sgdoptimizer alongwith three hidden layers and no BN & i.e. 0.68
6. With keras it become easy to implement deep neural network. While doing this assignment I I very powerful than classical modeling.

